Last chance to save

Prices rise for the Digiday Programmatic Marketing Summit after Mar. 24

REGISTER

DIGIDAY PUBLISHING SUMMIT

Inside The New York Times' AI newsroom strategy

By Sara Guaglione • September 22, 2025 •

We use cookies in compliance with the CCPA to make the website work better for you and to understand your needs. You can read more in our Privacy Policy.

Okay

This article is part of Digiday's coverage of its Digiday Publishing Summit.
More from the series →

The New York Times' editorial team is using AI technology to pursue a host of stories it couldn't tackle before, as they involved huge and messy datasets.

The team behind figuring out how to use AI to parse through hundreds of hours of video or thousands of datapoints is led by Zach Seward, The New York Times' editorial director of A.I. initiatives. His role was created in December 2023, part of a wave of new AI-focused positions that media companies formed to figure out which AI guidelines, projects and tools to develop for the newsroom to give reporters a competitive edge.

Onstage at the Digiday Publishing Summit in Miami, Fla., last week, Seward outlined how his team is working with reporters, what tools they've built and what the key use cases for AI are so far. Seward has a team of eight, including four engineers, a product designer and two editors.

Using AI for research and investigations is "by far the biggest use of our resources and I think the biggest opportunity right now when it comes to AI in media," Seward said. His team

mostly works by helping a reporter use AI technology for one project, and then creating a repeatable process from that experience for others in the newsroom to use.

Building the 'Cheat Sheet' AI tool

A Times reporter came to Seward's team with an impossible task – 500 hours of leaked Zoom recordings of an election interference group to go through before Election Day. AI tools were used to transcribe 5 million words and identify parts of those transcripts that were of interest to the reporter.

"[The election interference group] wasn't so dumb as to say, 'We're going to spread misinformation on the internet'... then you could Control F" to find that information in the transcripts, Seward said. "Where AI becomes useful – typically it's referred to as semantic search, or sometimes vibes-based searching – where you're looking for topics, concepts, things that are similar. And that's hugely useful when looking through enormous corpuses of text," Seward said. That led to a big story before the presidential election last year, he added.

Those efforts were then developed into a spreadsheet-based AI tool built internally called Cheat Sheet. Reporters can pick (with

TRENDING IN GENERATIVE AI

01 GENERATIVE AI

Butler/Till's first agentic media buying tests cut media and supply chain costs

02 MEMBER EXCLUSIVE

Future of Marketing Briefing: Agency operating systems face a differentiation problem

SPONSORED

How new infrastructure, like the Model Context Protocol, is reshaping marketing workflows

03 GENERATIVE AI

Shopify says purchases are coming 'inside ChatGPT' through agentic storefronts as OpenAI retreats on Instant Checkout

the guidance of Seward's team) which LLM model to use with Cheat Sheet, Seward said. It's now being used by several dozen reporters.

Stay up to date on Digiday events

Stay informed on who is attending, who is speaking and how you can save on passes

Business email

Job title

SIGN UP

Seward declined to share other specific AI tools The New York Times newsroom was using, though he said it was "pretty much all the commercial AI providers as well as open source models."

Cheat Sheet also helped a reporter who had an unorganized list of 10,000 names of people who had registered for a tax cut in Puerto Rico.

"You can't Google 10,000 names… but a computer can Google 10,000 names. And then using AI, we could analyze those search results for certain markers that [the reporter] was interested in," Seward said.

Even though the results weren't entirely accurate, it helped sort the names into more promising leads. The reporter could then call

them up and continue reporting out the story, Seward said.

Seward's team's approach is to took on an individual reporting challenge – "knotty, huge, messy data sets" with an "immediate deadline" – "but always with an eye toward building up tooling that will make that repeatable in the future," he said.

How Seward's team works with the newsroom

How does Seward's team decide which AI tools to build and where they can help the newsroom?

Seward said it comes down to constant communication with the newsroom. His team hosts training sessions on how to use AI tools for research investigations. Seward's team has spoken to 1,700 of the 2,000 people in the newsroom so far, he said.

The New York Times also has an open Slack channel that anyone from the newsroom can join to ask questions and share use cases – ranging from "how can I get Gemini?" to one bureau chief inspiring another across the world with an idea for how they're using AI technology.

"AI… is such a personal technology," Seward said. "The way people would describe what

they want out of AI can be different to the person." Many have experienced "writer's block in the chatbot… With a tool that can do anything for me, sometimes the challenge is… what can it do? And so we're just trying to help answer that question," he said.

Dealing with skepticism from the newsroom

AI isn't being used to write articles at The New York Times, Seward said. Reporters are allowed to use AI to draft copy around published articles, such as SEO and headlines, he said.

Seward said he reminds editorial staff to "never trust output from an LLM. Treat it… with the same suspicion you would a source you just met and you don't know if you could trust."

Some newsrooms have challenged upper management about using AI vendors to generate articles, or fought back when policies are rolled out that push for more use of the technology from editorial.

Seward team deals with any skepticism from reporters about using AI by acknowledging their reservations, and showing them how the technology can be useful, Seward said.

MOST READ

01 MEMBER EXCLUSIVE
Future of TV Briefing: How Netflix is shaking up the upfront deal model

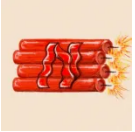

02 THE PROGRAMMATIC MARKETER
The Trade Desk remains the dominant DSP but its advertisers are starting to shop around

03 GENERATIVE AI
Butler/Till's first agentic media buying tests cut media and supply chain costs

04 MEMBER EXCLUSIVE
Future of Marketing Briefing: Agency operating systems face a differentiation problem

05 AI REVENUE GENERATION
Cloudflare's compliant crawler highlights tension – and opportunity – in the emerging AI content market

"We're not trying to be AI boosters. In fact, quite the opposite. I think there's a lot of caution. A lot of time we spend cautioning people about uses of AI, both [in the] legal and editorial senses," he said. "But if we can have you leave a session saying, 'I'm still pretty concerned about this whole environmental issue and maybe like destroying humanity thing – but in the meantime, it's going to let me transcribe handwritten notes in Arabic that I took messy iPhone photos of while I was on a reporting trip, and that's pretty cool.' And no reporter is going to say no to a competitive advantage, which I think is the theme of what we're trying to build for them."

What's the biggest challenge Seward faces in his newly-created role at The New York Times?

"I definitely live in fear of an error that is in some way attributable to AI. To be clear, we also say in sessions with our newsroom we would never attribute an error to AI, meaning it's always on us," Seward said.

"I would 100% feel responsible" if something like that happened, he added.

Share Copy

Sign up to get the Digiday Daily newsletter every morning

newsletter every morning

Business Email

Job title

SIGN UP

MORE IN MEDIA

VIEW MORE »

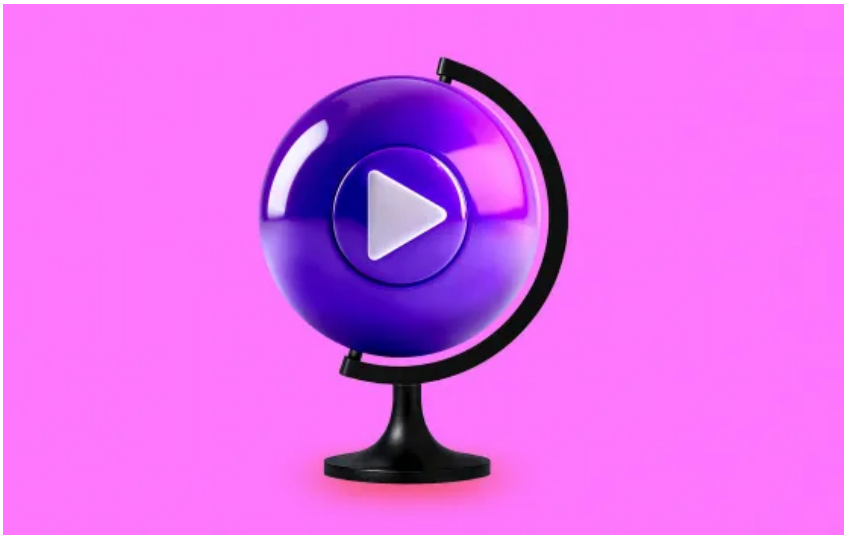

CONTENT & COMMERCE

Why Parker Thatch transformed its strip-mall storefront into a livestreaming studio

March 20, 2026

Parker Thatch recently remodeled its store to serve as a hybrid customer-facing retail experience and broadcast studio.

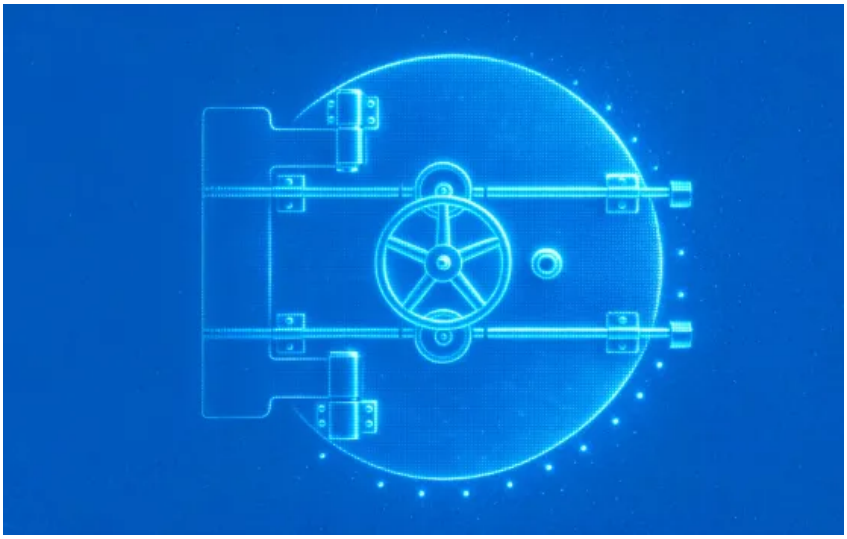

AI REVENUE GENERATION

Cloudflare's compliant crawler highlights tension – and opportunity – in the emerging AI content market

March 19, 2026

Cloudflare faces tension in its new role: sitting in the middle between publishers and AI companies while balancing trust, control and monetization.

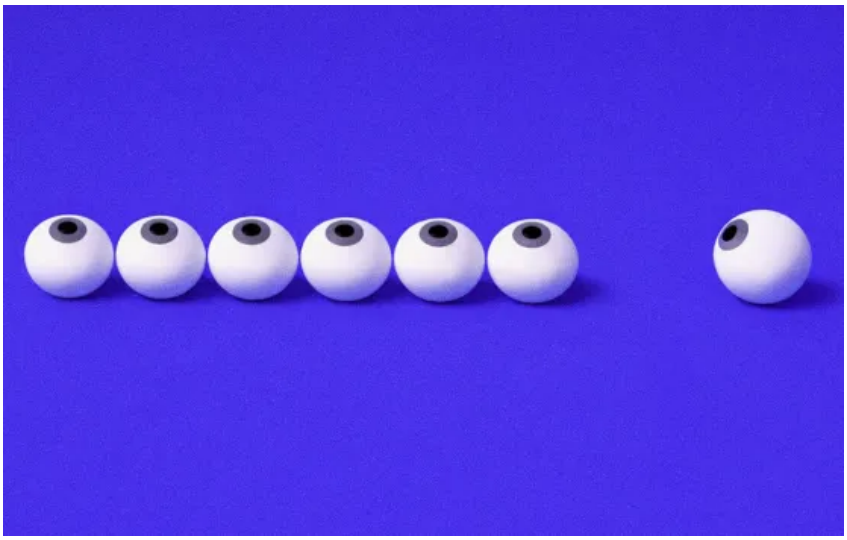

MEMBER EXCLUSIVE

Media Briefing: What to expect at the Digiday Publishing Summit, March 2026 edition

March 19, 2026

Execs from The Atlantic, Arena Group, Bloomberg, Business Insider, The Guardian, New York Post, People Inc., Washington Post, and more, will share their strategies on everything from zero-click audience strategy, to AI licensing deals and RAG readiness, to how they're embracing creator strategies to help boost engagement with younger audiences.

DIGIDAY+

Get access to tools and analysis to stay ahead of the trends transforming media and marketing

SUBSCRIBE

NEWSLETTER

Get Digiday's top stories every morning in your email inbox.

Business email

Job title

SIGN UP

CONNECT

Follow @Digiday for the latest news, insider access to events and more.

FAQ ADVERTISE PRIVACY POLICY MASTHEAD DIGIDAY MEDIA DIGIDAYMEDIA © 2026. ALL RIGHTS RESERVED